

When AI Shops with AI: Buyer-Seller Agent Negotiation and Market Outcomes in Agentic Commerce

Abstract

A growing share of online purchases will soon be negotiated not between a person and a merchant, but between a buyer’s shopping agent and a seller’s pricing agent, neither of which tires, forgets, or hesitates to walk away from a bad deal. This paper asks what happens to prices, welfare, and market structure when that shift occurs, using a matched agent-based computational economics (ACE) simulation of 105,600 transactions across 80 replicate markets, each run through five institutional regimes that span the space between today’s posted-price web store and a fully autonomous, multi-seller negotiation. Moving from limited human-style search to full AI-mediated search raises total welfare by improving buyer-seller matching, but pure price bargaining on top of that search is almost perfectly redistributive: it reallocates surplus between buyer and seller without enlarging the pie. Allowing agents to negotiate over bundles and delivery terms in addition to price, however, does enlarge it, by roughly 13.6 welfare units per transaction, a genuine efficiency gain from a richer negotiation space rather than a transfer. That efficiency gain has a darker mirror: as sellers’ agents learn more about a given buyer, the buyer’s share of the surplus falls from 84 percent to 68 percent, a direct and statistically precise price-discrimination effect. Full-market AI search also concentrates trade sharply among a small number of winning sellers (a win-share Herfindahl index rising from about 0.10 to 0.60), though allowing sellers to compete for the buyer’s attention simultaneously partially reverses this concentration. Two further findings speak directly to policy-relevant fears about agentic commerce: buyers who share a common buyer-side agent provider capture a modestly larger surplus share than inde-

pendent buyers, a real monopsony effect; and while shared seller-agent infrastructure does not mechanically compress market-wide price dispersion, it does explain roughly 45 percent of the markup variance among sellers on that infrastructure, evidence of real correlation without evidence of collusion. A smaller, exploratory pilot in which real Gemini and OpenRouter language models negotiate directly, rather than following the paper’s rule-based protocol, corroborates the central bundling result: conditional on reaching a deal, multi-attribute negotiation again outperforms price-only bargaining. Together, the results suggest that the economic consequences of agentic commerce depend less on whether AI agents are present than on which specific institutional rules govern how they search, bargain, and disclose information.

Keywords: agentic commerce, AI shopping agents, multi-agent negotiation, agent-based computational economics, algorithmic pricing, market design

Highlights

- AI-mediated search raises welfare; pure price bargaining only redistributes it
- Multi-attribute negotiation creates surplus: +13.6 welfare units versus price-only
- Seller information advantage cuts buyer surplus share from 84% to 68%
- Full-market search raises seller concentration sixfold; competition offsets it
- Shared buyer- and seller-agent infrastructure shifts bargaining power, not collusion

1. Introduction

Imagine a buyer who wants a new television. Twenty years ago she would have visited two or three stores, compared a handful of models, and either

paid the sticker price or haggled a little at the register. Today she opens a
5 handful of browser tabs and a comparison site does the work of visiting every
store at once. Tomorrow, on the trajectory that most large e-commerce and
AI platforms are currently building toward, she will simply tell a shopping
agent her budget and preferences and let it search, negotiate, and transact
on her behalf. On the other side of that transaction, an increasing number
10 of merchants already price and reprice using algorithmic pricing engines, and a
growing number are experimenting with agents that can negotiate directly with a
buyer’s agent rather than simply posting a take-it-or-leave-it price. Both trends
are old; what is new is the prospect that they meet in the same transaction, with
an AI agent representing the buyer negotiating against an AI agent representing
15 the seller, and no human present at the point of sale at all.

This paper asks a simple question with a genuinely open answer: does it
matter, for prices, welfare, and market structure, precisely how that meeting is
organized? The intuition that "AI agents will change e-commerce" is by now
close to a truism, echoed across the trade press and the emerging academic
20 literature on agentic markets (Bichler, 2026; Shahidi et al., 2025; Hadfield and
Koh, 2025). It is a much less obvious claim that the institutional details of the
meeting – how much of the market a buyer’s agent actually searches, whether
the resulting price is posted or bargained, whether the bargaining space includes
only price or also bundles and delivery terms, and whether the buyer’s or seller’s
25 agent is one of many independent implementations or one of a handful of shared
platforms serving many principals at once – determine whether that change
helps buyers, helps sellers, or mostly just reshuffles who captures a roughly
fixed amount of surplus. Untangling those institutional details from the mere
presence of AI is the central contribution of this paper.

30 The question has an obvious empirical difficulty: agentic commerce, in the
sense of autonomous buyer and seller agents transacting without a human in
the loop, is not yet common enough to observe at scale, and the platforms
that are beginning to deploy something like it treat their negotiation logic as
a trade secret. Two research strategies are available in the interim. One is to

35 build small experimental testbeds and have real large language models negotiate
 inside them, which has the appeal of studying the actual technology but the
 limitation that current testbeds are necessarily small, and the artifacts of any
 one model’s training and prompting can be mistaken for a general economic
 finding (Bianchi et al., 2024; Zhu et al., 2025; Bansal et al., 2025). The other
 40 is to build an explicit, parametrized model of the market mechanism itself – an
 agent-based computational economics (ACE) approach in the tradition surveyed
 by Tesfatsion (2006) – and ask what that mechanism implies for prices and
 welfare under full experimental control. This paper takes the second route
 as its primary empirical strategy, and supplements it with a small, explicitly
 45 exploratory pilot of the first kind, run over the same underlying environment,
 as a check on whether the mechanism-level story survives contact with an actual,
 non-scripted negotiator.

The simulation crosses five institutional regimes – posted price with limited
 search, AI-assisted full-market search with no bargaining, bilateral price-only
 50 bargaining, bilateral multi-attribute bargaining, and competitive multi-seller
 solicitation – with seller information about the buyer, whether the negotiation
 space includes non-price attributes, and whether either side of the market is
 served by a small number of shared agent providers or a large number of inde-
 pendent ones. Every replicate market is run through every applicable combina-
 55 tion of these conditions with the same drawn population of buyers and sellers, so
 that differences across regimes reflect the institutional rule change itself rather
 than a different draw of who happened to be trading with whom. The result
 is 105,600 simulated transactions across 80 replicate markets, generating the
 six empirical findings summarized in the abstract: search improves matching
 60 efficiency; price-only bargaining is close to purely redistributive; multi-attribute
 bargaining genuinely creates surplus; richer seller information about the buyer
 extracts surplus from her; full AI-mediated search concentrates trade among
 fewer winning sellers, an effect that competitive multi-seller solicitation par-
 tially undoes; and buyer-side and seller-side agent-provider concentration both
 65 leave detectable, but importantly different, fingerprints on the resulting prices.

The paper’s contribution is therefore not to claim that AI agents will make markets more efficient or less fair in general, a claim the results below suggest is not well posed, but to show which specific design choices move each of those two outcomes, and by how much, inside one consistent and fully reproducible
70 environment. Section 2 situates the paper against four adjacent literatures: e-commerce search and recommendation, the nascent literature on agentic markets, algorithmic pricing and collusion, and the older literatures on price discrimination and bilateral negotiation that this new setting recombines. Section 3 develops the formal model, the five institutional regimes, and seven hypothe-
75 ses that map directly onto the empirical tests that follow. Section 4 describes the simulation and the supplementary LLM-agent pilot, and is deliberately explicit about the difference in what each can and cannot claim to show. Section 5 reports the results, Section 6 discusses their implications for platform design and competition policy, and Section 7 concludes with the paper’s limitations
80 and the empirical work that would be needed to move from a mechanism-level claim to a claim about any specific deployed system.

2. Literature Review

2.1. *Search costs and matching in electronic markets*

The starting premise of this paper – that reducing the cost of search should,
85 in itself, be welfare-improving – rests on a literature that predates AI shopping agents by decades. Stigler (1961) first formalized the idea that costly search generates equilibrium price dispersion even among sellers of an identical good, because no buyer can profitably examine every seller. Bakos (1997) showed that electronic marketplaces, by collapsing the marginal cost of examining one more
90 seller toward zero, should compress that dispersion and shift surplus toward buyers, a prediction that generated a lively empirical debate about whether real online markets actually behaved that way; Harrington (2001)’s comment on Bakos’s model is a useful reminder that the mapping from "search is cheaper" to "buyers do better" depends on exactly how search frictions are modeled, a

95 caution this paper takes seriously by treating search breadth as one dial among
several rather than the only one that matters. More recent work on AI-mediated
recommendation and choice, such as Kim et al. (2023)'s study of choice overload
in conversational recommendation, extends the search-cost logic into the current
generation of AI-assisted shopping tools, and is one of the few empirical anchors
100 available for how an AI intermediary changes buyer decision quality rather than
just buyer effort.

This paper's regimes M0 and M1 are a direct, if stylized, operationalization
of that debate: M0 restricts the buyer's agent to examining a small, fixed num-
ber of sellers, mimicking a human shopper's practical search limit, while M1
105 gives the buyer's agent costless access to the entire market. The resulting wel-
fare gain from M0 to M1, and the accompanying fall in the average transaction
price, replicate the classic search-cost story quantitatively rather than just qual-
itatively, and set up the paper's more novel question: once search is essentially
free, as it plausibly becomes once an AI agent is doing it, does adding bargaining
110 on top of that search create any further gains, or does it only redistribute the
ones search already delivered?

2.2. *Agentic commerce and markets with AI agents*

A body of work published almost entirely within the last two years has be-
gun to ask what happens when the buyer, the seller, or both are themselves
115 autonomous software agents rather than a human using a tool. Bichler (2026)
frames the central open question of this literature precisely: that agentic mar-
kets could in principle deliver large efficiency gains from better search and faster
negotiation, but that "gains are not automatic" and depend on the market's in-
stitutional design, a framing this paper takes as its organizing hypothesis and
its title's implicit claim. Shahidi et al. (2025) develop a formal argument, which
120 they term the Coasean Singularity, that falling transaction costs from AI in-
termediation should push real markets toward the frictionless, transfer-neutral
benchmark that Coase's theorem describes in the limit – a benchmark against
which this paper's finding that price-only bargaining is nearly perfectly redis-

125 tributive is a striking, if partial, confirmation. Hadfield and Koh (2025) take a
 broader view of an emerging "economy of AI agents" and the institutional infras-
 tructure, from identity and liability to contracting norms, that such an economy
 will require. Closer to this paper's specific empirical object, Allouah et al. (2026)
 directly study what a shopping AI agent actually buys and how its choices are
 130 biased relative to a human's, an important complement to a paper like this one
 that instead holds buyer preferences fixed and varies the market's rules. Bansal
 et al. (2025) and Zhu et al. (2025) each build open experimental testbeds –
 Magentic Marketplace and an agent-to-agent consumer negotiation benchmark,
 respectively – for studying real deployed language models transacting with each
 135 other, the same broad research strategy this paper's supplementary LLM pilot
 adopts at a far smaller scale. Bichler et al. (2025) model the game-theoretic
 dynamics of learning agents in agentic markets directly, a formally richer com-
 plement to this paper's rule-based approach. Bianchi et al. (2024) is an earlier
 and by now foundational benchmark for evaluating how well large language mod-
 140 els negotiate against each other or against scripted opponents, and its finding
 that negotiation outcomes are highly sensitive to which specific model is playing
 which role is directly relevant to this paper's own small-scale finding, in Section
 5.6, that a Gemini-versus-Gemini pairing, an OpenRouter-versus-OpenRouter
 pairing, and a mixed pairing do not behave identically. Finally, Bostoen and
 145 Krämer (2026) examine how well existing digital-market regulation, specifically
 the European Union's Digital Markets Act, anticipates a world of autonomous
 negotiating agents, a question this paper's concentration findings speak to di-
 rectly in Section 6.

What is missing from this fast-moving literature, and what this paper tries
 150 to supply, is a controlled comparison across institutional regimes using a single
 population of buyers and sellers held fixed across treatments. Papers that build
 testbeds for real language models are, almost by construction, exploring one or a
 few configurations of that testbed rather than sweeping systematically across the
 space of market designs; this paper inverts that trade-off, sacrificing the realism
 155 of an actual deployed language model as the primary evidence in exchange for

full experimental control over exactly which institutional lever is being turned.

2.3. Algorithmic pricing and the collusion question

A parallel and more mature literature asks whether algorithmic pricing agents, even without any agentic buyer on the other side, can produce collusive-like outcomes without any explicit agreement, human intent, or even communication among the pricing algorithms. Calvano et al. (2020) showed in a landmark reinforcement-learning simulation that independently trained pricing algorithms could learn to sustain supra-competitive prices through reward-punishment schemes that require no explicit signaling, a result extended to settings with imperfect monitoring in Calvano et al. (2021). Klein (2021) demonstrated a similar tacit-collusion result for Q-learning agents under sequential pricing, and Brown and MacKay (2023) brought some of this theoretical concern to real-world grounding with evidence on how actual algorithmic pricing software affects retail gasoline competition. Sanchez-Cartas and Katsamakas (2025) extend this line of inquiry to platform competition specifically, which is closer to this paper’s e-commerce setting.

This paper’s own concentration findings sit adjacent to, but are conceptually distinct from, that collusion literature. The algorithmic-collusion papers ask whether independent pricing agents can converge to a joint-profit-maximizing outcome without communicating; this paper instead asks what happens when sellers are not fully independent to begin with, because a subset of them share a common third-party agent provider, and finds a genuinely nuanced answer, developed fully in Section 5.5, in which shared infrastructure correlates within-network markups without mechanically compressing overall market-wide price dispersion. That distinction matters for how the result should be read: it is evidence of correlated behavior among agents on shared infrastructure, not evidence that agentic commerce is inherently collusive, and Section 6 argues that conflating the two would overstate what the simulation shows.

2.4. *Personalization and price discrimination*

185 The seller-information dimension of this paper’s design draws on an established literature on how much a firm can extract from a buyer once it can infer that buyer’s willingness to pay. Bergemann et al. (2015) characterize the full range of consumer-surplus and producer-surplus outcomes that price discrimination can theoretically produce as a function of how much the seller knows
190 about buyer valuations, a range this paper’s own information-level treatment – from a seller who knows nothing beyond its own cost to one that observes the buyer’s urgency and purchase history – is designed to traverse empirically rather than just characterize theoretically. Dubé and Misra (2023) bring this question to real personalized-pricing data and find that welfare effects of personalization depend heavily on the specific pricing technology in use, a caution
195 against assuming personalization is uniformly good or bad for consumers. Acquisti and Varian (2005) study the specific mechanism of conditioning price on purchase history, directly analogous to this paper’s "history-aware" seller information level, and find, as this paper does, that loyalty signals can be turned
200 against the buyer rather than rewarded. Akerlof (1970)’s classic analysis of information asymmetry, though framed around quality rather than willingness to pay, is the conceptual ancestor of this entire literature: whichever side of a transaction knows more has a structural advantage, and an AI agent’s central economic function is to change who knows what.

205 2.5. *Negotiation theory and transaction costs*

 Finally, the paper’s bilateral bargaining regimes rest on a negotiation-theoretic foundation reaching back to Nash (1950)’s axiomatic characterization of the bargaining outcome and Rubinstein (1982)’s non-cooperative alternating-offers model, which showed that patience and the cost of delay, not raw bargaining
210 power alone, determine how a fixed surplus gets split between two parties. The simulation’s actual negotiation protocol is a parametrized, monotonic-concession model in the tradition of Faratin et al. (1998), whose negotiation decision functions remain the standard reference for modeling how an autonomous agent

should concede over time as a function of its own urgency and its beliefs about
 215 its counterpart. Williamson (1981) and Coase (1937) complete the theoretical
 picture from the institutional-economics side: both argue that the boundaries
 of firms and the design of contracts exist to economize on the transaction costs
 of negotiating, monitoring, and enforcing exchange, which is precisely the cost
 category that an AI negotiating agent is built to reduce. Taken together, this
 220 literature offers a sharp prediction that this paper’s own results confirm empir-
 ically rather than merely restate: because a bilateral bargain over a fixed pie is,
 in Rubinstein’s framework, a transfer problem, a bargaining regime layered on
 top of an already-efficient search regime should not be expected to raise total
 welfare on its own. Any welfare gain from negotiation, if the theory is right,
 225 has to come from expanding what is actually being bargained over, which is
 exactly the test that distinguishes this paper’s price-only regime (M2) from its
 multi-attribute regime (M3).

3. Model and Hypotheses

3.1. Buyers, sellers, and utility

230 Each simulated buyer i has a private valuation v_i for one unit of a homoge-
 neous product, drawn independently across buyers, and a budget ceiling slightly
 above that valuation. Each seller j has a private unit cost c_j and a minimum ac-
 ceptable markup over that cost, both drawn independently across sellers. When
 a transaction is completed at price p , and, in regimes that allow it, with a bun-
 235 dle add-on of realized value b to the buyer at incremental cost b_c to the seller,
 the buyer’s realized surplus and the seller’s realized surplus are

$$CS = v_i + b - p, \quad (1)$$

$$PS = p - c_j - b_c, \quad (2)$$

so that total welfare from the transaction is $W = CS + PS = v_i + b - c_j - b_c$,
 which is maximized, holding the identity of the trading pair fixed, by realizing
 the bundle whenever its value to the buyer exceeds its cost to the seller ($b > b_c$),

240 and is otherwise invariant to the price p at which the transaction clears. This is
the model’s formal statement of the redistribution logic developed in Section 2.5:
for a fixed buyer-seller pair and a fixed bundle decision, the price only reallocates
 W between CS and PS , and does not change W itself. Total welfare can still
change across institutional regimes, but only through two channels that this
245 decomposition makes explicit: which buyer ends up matched to which seller (a
search or matching effect), and whether a value-creating bundle is realized at
all (a bargaining-scope effect).

3.2. Five institutional regimes

The five regimes vary how a buyer’s agent finds a seller and whether, and
250 over what, the two agents then bargain.

- **M0 (posted price, limited search).** The buyer’s agent examines two
of sixteen sellers at random and pays the lower of their two posted prices,
mimicking the practical search limit of a human shopper who visits a small
number of stores.
- 255 • **M1 (AI-assisted full-market search).** The buyer’s agent costlessly
compares all sixteen sellers’ posted prices and buys from the cheapest,
with no bargaining.
- **M2 (bilateral price-only bargaining).** The buyer’s agent identifies
the seller that full-market search would have selected, then negotiates a
260 price with that single seller under a monotonic-concession protocol; the
bargaining space is restricted to price alone.
- **M3 (bilateral multi-attribute bargaining).** Identical to M2, except
that the seller’s agent may, when it is profitable to do so, offer a bundle
add-on – an accessory, an extended warranty, faster delivery – in place of,
265 or alongside, a pure price concession.
- **M4 (competitive multi-seller solicitation).** The buyer’s agent solic-
its simultaneous offers from the cheapest seller identified by search plus

two additional randomly encountered rivals, discloses the best rival offer
obtained so far to each subsequent seller, and accepts whichever result-
270 ing offer maximizes the buyer’s realized utility. M4 nests M2/M3 as the
special case of a single solicited seller, so that any advantage of M4 is
attributable to competitive pressure and disclosure rather than to a sys-
tematically luckier draw of counterparties.

Every regime is additionally crossed with four levels of seller information
275 about the buyer (low, preference-aware, history-aware, and high, the last com-
bining urgency and history signals into a stronger inference of the buyer’s
willingness to pay), with whether the bargaining space is price-only or multi-
attribute, with whether sellers are independent or clustered on a small number
of shared third-party pricing-agent providers, and with whether buyers are in-
280 dependent or pooled on a common buyer-agent provider that can aggregate
their bargaining power. Every replicate market draws its own population of
sellers and buyers once, and then runs the *same* population of buyers through
every regime and condition combination applicable to that replicate, so that the
comparisons reported in Section 5 are within-market rather than across freshly
285 redrawn populations.

3.3. Hypotheses

The five regimes and four crossed conditions generate seven hypotheses, each
of which is tested directly against one of the results reported in Section 5.

H1 (search improves matching). Moving from limited search (M0) to full-
290 market AI-mediated search (M1) raises total welfare, because a buyer’s
agent that examines the whole market is more likely to match with a
seller whose cost or fit is genuinely favorable rather than merely favorable
among the two sellers a human happened to visit.

H2 (bargaining is redistributive). Adding price-only bilateral bargaining
295 on top of full-market search (M2 relative to M1) does not raise total wel-

fare, because the identity of the matched buyer-seller pair is unchanged and the bargaining space contains nothing but price.

H3 (multi-attribute bargaining creates surplus). Allowing the bargaining space to include a bundle add-on (M3 relative to M2) raises total welfare whenever that bundle is worth more to the buyer than it costs the seller to provide, an efficiency gain that price-only bargaining cannot capture.

H4 (competitive solicitation moderates concentration). Soliciting multiple sellers simultaneously and disclosing rival offers (M4) reduces the concentration of trade among a small number of winning sellers, relative to the single-seller search-then-bargain regimes (M1 through M3).

H5 (information enables extraction). As the seller’s information about the buyer’s valuation and urgency increases, the buyer’s share of the transaction surplus falls, because a better-informed seller can price closer to the buyer’s true reservation value.

H6 (buyer-side concentration confers bargaining power). Buyers served by a common, high-volume buyer-agent provider capture a larger share of the transaction surplus than buyers served by independent agents, reflecting the aggregated bargaining leverage such a provider can bring to bear even when negotiating on behalf of any one individual buyer.

H7 (seller-agent infrastructure correlates prices without necessarily concentrating them). Sellers who share a common third-party pricing-agent provider show more correlated markups with one another than independent sellers do, but this within-network correlation does not necessarily translate into lower market-wide price dispersion, since a market containing a few large, distinct provider networks can be just as heterogeneous overall as one containing many fully independent sellers.

4. Methodology

4.1. *The rule-based ACE simulation as the paper’s primary evidence*

325 The results reported in Section 5 are generated entirely by a rule-based agent-based computational economics simulation, not by any large language model. This is a deliberate methodological choice, not a limitation to be apologized for. The simulation’s negotiation protocol – a parametrized, monotonic-concession function of each agent’s own urgency, its inferred belief about its counterpart’s
330 reservation value, and, where applicable, a competitor’s disclosed offer – follows the negotiation-decision-function tradition of Faratin et al. (1998) and is implemented deterministically given a fixed random seed, so that every reported number in this paper can be regenerated exactly from the published code. That determinism buys three things that no study of an actual deployed AI system
335 can currently offer: full experimental control over the primitives that matter for the theory (buyer valuations, seller costs, and each side’s information about the other), a within-market matched design in which the same buyers are run through every regime, and complete reproducibility. The cost of this choice is equally direct: the simulation’s negotiation behavior is calibrated qualitatively
340 to the mechanisms described in the agentic-commerce literature reviewed in Section 2.2, not fit to any particular deployed language model’s actual behavior, and the numbers in Section 5 should accordingly be read as properties of an explicit, transparent model of the institutional mechanism, not as a forecast of what any specific commercial AI shopping agent will do.

345 Sellers’ costs are drawn independently and uniformly, and their minimum acceptable markups and bundle values are drawn independently as well; sellers on a shared third-party provider retain their own independently drawn costs but have their strategic parameters (minimum margin, in particular) drawn from one of two shared strategy clusters, capturing the idea that a common pricing-agent
350 vendor imposes a common pricing policy on its client sellers without literally giving them a common cost structure. Buyers’ valuations and budgets are likewise drawn independently and uniformly, and a small, fixed idiosyncratic taste term

is added to each buyer’s ranking of otherwise similar sellers; without that term, price-ranked full-market search collapses into a knife-edge Bertrand outcome in which the single cheapest seller wins essentially every buyer in the market, an artifact of assuming products are perfectly undifferentiated that the idiosyncratic term relaxes into a strong, but not total, winner-take-all tendency, which is both more realistic and closer to what the resulting Herfindahl indices in Section 5.3 actually show. The buyer-side monopsony mechanism is implemented at the level of the individual buyer, not the market: within a market flagged as served by a common buyer-agent provider, only the roughly seventy percent of buyers who are themselves drawn as clients of that provider receive the associated bargaining-power adjustment, so that the aggregate effect reported in Table 4 reflects a genuinely heterogeneous population of buyers rather than a uniform market-wide policy switch. The resulting dataset comprises 105,600 simulated transactions across 80 replicate markets, and every table and regression in Section 5 is computed with standard errors clustered at the replicate-market level to account for the fact that transactions drawn from the same market are not independent observations.

4.2. *A supplementary pilot with real language models*

Because the simulation above does not call any large language model, a natural and important question is whether its central qualitative predictions survive contact with an actual, non-scripted negotiator. To address this without overstating what a small pilot can show, a second, much smaller experiment has two real language models – Google’s Gemini (the `gemini-flash-lite-latest` tier) and a rotating set of free-tier models accessed through OpenRouter – play the buyer and seller roles directly. Each agent receives its own private valuation or cost information in its system prompt, is asked to negotiate over a strict JSON action protocol (propose a price and, where applicable, a bundle; accept the counterpart’s last proposal; or walk away), and negotiates for up to six rounds. The pilot crosses a price-only regime and a multi-attribute regime, analogous to M2 and M3, with three model pairings – Gemini negotiating against Gemini,

OpenRouter against OpenRouter, and a mixed pairing with Gemini as the buyer and OpenRouter as the seller – across six paired replicate draws from the same underlying valuation, cost, and bundle-value distributions used in the main simulation, for 36 negotiation episodes in total.

This pilot is reported throughout as an exploratory robustness check, not as the paper’s evidentiary core, for reasons that are worth stating plainly rather than burying in a footnote. Thirty-six episodes is not a powered experiment by the standards this paper otherwise holds itself to, and real free-tier language model calls occasionally fail to return a parseable action at all: six of the thirty-six episodes here ended in an unparseable or malformed response rather than a negotiated outcome or an explicit walk-away, and every one of those failures is logged and treated as a zero-welfare non-transaction rather than silently discarded. Because those failures happened to land more often in the multi-attribute cell than the price-only cell, an unconditional comparison of average welfare across the two regimes is confounded: multi-attribute negotiation looks worse only because more of its episodes failed to close, not because successful multi-attribute deals generate less surplus. Section 5.6 therefore reports both the unconditional comparison and a deal-conditional comparison that isolates surplus creation among episodes that actually reached an agreement, which is the comparison that speaks to the same mechanism as the M2-versus-M3 test in the main simulation.

5. Results

5.1. Search, bargaining, and the limits of pure price negotiation

Table 1 reports average price, consumer surplus, seller surplus, and total welfare in each of the five regimes, pooled across all information levels, bargaining-scope conditions, and concentration conditions. Total welfare rises from 416.1 under limited human-style search (M0) to 442.8 once the buyer’s agent can search the full market at zero cost (M1), a gain fully consistent with H1: a buyer’s agent that compares all sixteen sellers rather than two is simply more

likely to end up matched with a seller who offers her genuinely good value, independent of any bargaining. The average price also falls sharply over the same move, from 663.1 to 614.0, the classic search-cost-compression result of Bakos
415 (1997) reproduced quantitatively inside this paper’s own environment.

What happens next is the paper’s first genuinely informative negative result. Layering bilateral price-only bargaining on top of that same full-market search (M2) leaves total welfare completely unchanged at 442.8, to three significant figures identical to M1, even though the average price rises slightly, to 637.5,
420 and both consumer surplus and seller surplus shift relative to the no-bargaining outcome. This is exactly what H2 predicts and exactly what the negotiation-theoretic logic of Section 2.5 requires: because the matched buyer-seller pair is the same one full-market search would have selected regardless, and because the bargaining space in M2 contains nothing but price, the negotiation can
425 only move surplus between the buyer and the seller, not create or destroy it. A regression of total welfare on regime, bargaining scope, and concentration condition, reported in Table 7, confirms this numerically: the coefficients on the M1 and M2 regime indicators, relative to M0, are statistically indistinguishable from one another (26.67 versus 26.67, both $p < 0.001$), while the coefficient on
430 price-only bargaining scope (relative to multi-attribute) is a precisely estimated -13.67 ($p < 0.001$) – the same magnitude, with the opposite sign, as the welfare gain reported next.

That symmetry is the paper’s central positive result. Allowing the same bilateral negotiation to range over a bundle add-on as well as price (M3) raises
435 total welfare to 456.4, a gain of 13.6 welfare units over M2 (Table 1; see also `headline_numbers.json`), consistent with H3. Table 2 decomposes this gain: within M3, transactions in the multi-attribute condition average 456.4 in welfare against 442.8 in the price-only condition run over the identical population, and the bundle is actually realized in 42 percent of multi-attribute M3 trans-
440 actions (rising to 52 percent under the additionally competitive M4 regime), with an average net surplus of \$32.8 when it is realized in M3 and \$38.2 in M4. Figure 1 plots the resulting stacked consumer- and producer-surplus decompo-

sition across all five regimes, and Figure 3 isolates the price-only-versus-multi-attribute comparison within M3 and M4 directly. The competitive regime M4
445 delivers a similar multi-attribute welfare level, 454.9, essentially statistically indistinguishable from M3’s 456.4, which is itself informative: once the bargaining space already includes bundling, adding competitive multi-seller pressure on top does not meaningfully enlarge the pie further, though Section 5.3 shows it does something else important to how trade is distributed across sellers.

450 5.2. Seller information and the extraction of consumer surplus

Table 3 reports average price and each side’s surplus share as the seller’s information about the buyer rises from low, to preference-aware, to history-aware, to high (a combination of urgency and history signals). The consumer surplus share is essentially flat across the first three levels, at 0.84, 0.84, and 0.83
455 respectively, and then drops sharply to 0.68 once the seller has high information, with the average price rising correspondingly from about 615 to 696. Table 7’s information regression makes the same point with precision: relative to a low-information seller, the coefficient on high seller information is -0.160 on consumer surplus share ($p < 0.001$, cluster-robust standard errors), an order of
460 magnitude larger than the preference-aware coefficient of -0.001 or the history-aware coefficient of -0.009 . Figure 2 plots this relationship directly.

This pattern is consistent with H5 and with the broader price-discrimination literature reviewed in Section 2.4 (Bergemann et al., 2015; Acquisti and Varian, 2005): information about a buyer is not dangerous to that buyer in small doses
465 – knowing a stated preference or an isolated purchase history barely moves the split of surplus – but becomes dangerous once it is combined into a genuine estimate of urgency and willingness to pay, at which point a seller’s agent can price close enough to the buyer’s true reservation value that the buyer captures barely two-thirds of the value she places on the good. The practical implication, devel-
470 oped further in Section 6, is that the risk agentic commerce poses to consumer welfare through personalization is concentrated in a specific kind of information – inferred urgency and behavioral history combined, not any single data point

in isolation – which is a more actionable target for disclosure or platform-design intervention than "AI personalization" treated as an undifferentiated category.

475 *5.3. Search-driven concentration and its competitive antidote*

Table 5 (Herfindahl index panel) and Figure 4 report the Herfindahl-Hirschman index of seller win-shares – the standard concentration measure, here computed over which seller wins each transaction – across regimes. Under limited human-style search (M0), the HHI is close to its floor at about 0.10, since a buyer who
480 examines only two of sixteen sellers essentially spreads trade close to evenly. The moment search becomes costless and market-wide (M1 through M3), the HHI roughly quintuples to about 0.60 among independent sellers, since every buyer’s agent is now comparing the same full set of sellers and systematically routing trade toward whichever one offers the best combination of price and
485 idiosyncratic fit. This is the paper’s clearest illustration of a genuine tension inside H1’s welfare gain: the same full-market search that improves matching efficiency also concentrates trade sharply, a textbook instance of efficient markets not being the same thing as competitive market structure.

Consistent with H4, the competitive regime M4 – in which the buyer’s agent
490 solicits several sellers simultaneously and discloses the best rival offer to each – brings the HHI back down to about 0.45, a meaningful but partial reversal of the search-driven concentration documented above. The reversal is partial rather than complete because M4 still nests the same underlying idiosyncratic-fit and cost heterogeneity that drives concentration under M1 through M3; simulta-
495 neous solicitation redistributes some trade toward sellers who would otherwise have lost the buyer entirely, but it does not eliminate the advantage of being genuinely the best-suited seller for a given buyer. The practical reading is that competitive multi-seller solicitation is a real, if incomplete, structural remedy for the concentration that AI-mediated search would otherwise produce on its
500 own, a point with direct implications for how e-commerce platforms should design the default number of sellers a shopping agent solicits before committing to one.

5.4. Buyer-side concentration: a real, if modest, monopsony effect

Table 4 compares consumer surplus share for buyers served by a common, high-volume buyer-agent provider against buyers served by independent agents. Buyers on the common provider capture a consumer surplus share of 0.80, against 0.79 for independent buyers, and Table 7’s monopsony regression confirms this is not noise: the coefficient on being an independent buyer, relative to being on the common provider, is -0.016 ($p = 0.004$, cluster-robust standard errors), holding regime and seller information constant. Because this effect is implemented at the level of the individual buyer – only buyers actually flagged as clients of the common provider receive the associated bargaining-power adjustment, not every buyer in a market where such a provider happens to be present – the resulting aggregate gap is a genuine composition effect of roughly seventy percent of buyers in that condition receiving a real discount, not an artifact of a market-wide policy switch, which is precisely what H6 predicts and no more than that.

The magnitude here matters as much as its statistical significance. A one-and-a-half percentage-point shift in surplus share is a real transfer, consistent with the idea that pooling many buyers’ negotiating power onto one agent provider generates genuine aggregated bargaining leverage, in the spirit of a monopsony buyer using its scale to extract better terms. But it is a modest transfer relative to the far larger effects documented in Sections 5.1 and 5.2, and the paper does not find evidence that buyer-side agent concentration is, on its own, a first-order determinant of who captures the gains from agentic commerce. Whether it becomes one at greater scale – if, for instance, a small number of buyer-agent platforms eventually intermediate a large majority of retail purchases – is a natural and important extension that this paper’s simulation is well suited to test, but does not itself address at the population scale examined here.

5.5. Seller-agent infrastructure: correlated prices without collusion

Table 6 compares the standard deviation of sellers' markups and the coefficient of variation of transaction prices between markets where sellers are fully independent and markets where a subset of sellers share a common third-party pricing-agent provider. The result runs against the naive intuition that shared infrastructure should straightforwardly compress or homogenize prices: markup dispersion is if anything slightly *higher* under shared infrastructure (0.0737 versus 0.0684), and the price coefficient of variation is likewise marginally higher (0.0667 versus 0.0600). Table 7's price regression adds a second, complementary fact: sellers on shared infrastructure charge a lower average price, by 11.79 units ($p < 0.001$), than independent sellers, controlling for regime and information level – shared-provider sellers are, on average, somewhat more aggressive on price, not less.

At the same time, membership in a common provider's client network explains approximately 45 percent of the cross-seller variance in average markup among sellers who share that network, evidence of real, substantial within-network price correlation. Reading these two facts together yields the paper's most carefully qualified finding, and the one Section 6 argues requires the most care in how it is communicated: a market containing a few large but internally correlated seller networks can display just as much, or even slightly more, aggregate price heterogeneity as a market of many fully independent sellers, because the relevant comparison is not "correlated versus uncorrelated" but "a few large clusters versus many small ones," and clusters can be internally similar while still differing substantially from one another. This is consistent with H7 as stated, and is deliberately not framed as evidence that AI pricing agents collude: nothing in this simulation involves communication, signaling, or a repeated-game punishment mechanism of the kind that generates genuine tacit collusion in Calvano et al. (2020) or Klein (2021). What the simulation shows is a mechanical consequence of shared strategy parameters within a client network, which is a real phenomenon regulators and researchers should be able to distinguish from collusion, not a demonstration that the two are the same

thing.

5.6. A supplementary check: real language models negotiating

Table 8 reports the results of the 36-episode LLM-agent pilot described in
565 Section 4.2. Overall, 30 of the 36 episodes reached a negotiated agreement;
the remaining six ended in an unparseable model response, logged as a failed
episode rather than discarded. Deal rates varied noticeably by model pairing and
regime: the Gemini-versus-Gemini pairing closed every single episode in both
regimes, while the OpenRouter-versus-OpenRouter pairing closed only half of its
570 multi-attribute episodes (three of six), evidently because the added complexity
of reasoning jointly about a price and a bundle decision, on top of an already
demanding JSON output format, strains some free-tier models more than a
price-only negotiation does. This is itself a modest but genuine piece of evidence,
echoing Bianchi et al. (2024)’s finding that negotiation competence is highly
575 model-dependent, that the paper’s institutional-design conclusions should not
be read as claims about how any specific commercial AI agent will behave in
practice.

The comparison that matters for this paper’s central claim, however, is the
deal-conditional one. Averaged across all three model pairings, and conditioning
580 on episodes that actually reached an agreement, mean welfare in the multi-
attribute regime is 542.8 against 512.4 in the price-only regime, a gain of 30.4
welfare units. This is the same sign, and a larger magnitude, as the 13.6-
unit gain the main rule-based simulation attributes to the identical mechanism
in Section 5.1. An unconditional comparison across all episodes, by contrast,
585 shows the multi-attribute regime performing *worse* on average, purely because
more of its episodes failed to close for the reasons just described; reporting only
that unconditional number would have created a spurious contradiction with the
paper’s central finding, which is why both numbers are reported here and the
deal-conditional comparison, not the unconditional one, is the basis for treating
590 this pilot as corroborating evidence. Thirty-six episodes across three model
pairings is far too small a sample to be the paper’s evidentiary core, and it is

not offered as one; it is offered as a check that the efficiency logic documented in the rule-based simulation is not simply an artifact of that simulation’s own functional-form assumptions, and on that narrow question it passes.

Table 1: Average outcomes by institutional regime, pooled across all information, bargaining-scope, and concentration conditions.

Regime	Price	Consumer surplus	Seller surplus	Welfare	Rounds
M0 (posted, limited search)	663.1	316.2	100.0	416.1	0.0
M1 (full-market search)	614.0	365.3	77.5	442.8	0.0
M2 (bilateral, price-only)	637.5	341.8	101.0	442.8	3.5
M3 (bilateral, multi-attribute)	637.5	361.1	88.5	449.6	3.5
M4 (competitive, multi-attribute)	639.2	367.5	80.6	448.0	3.4

Table 2: Welfare decomposition by bargaining scope within the multi-attribute regimes.

Regime	Scope	Welfare	Bundle realized	Net bundle surplus
M2	price-only	442.8	—	—
M3	price-only	442.8	0%	—
M3	multi-attribute	456.4	42%	32.8
M4	price-only	441.2	0%	—
M4	multi-attribute	454.9	52%	38.2

595 6. Discussion

The six results in Section 5 add up to a more specific claim than "AI agents will change e-commerce," and it is worth stating that claim in the terms Bichler (2026) uses to frame the same question: gains from agentic markets are not automatic, and this paper’s contribution is to locate precisely where, inside a single
600 consistent environment, they are automatic and where they are not. Search is close to automatically welfare-improving, because giving a buyer’s agent access to the whole market rather than a small sample of it is close to a pure information improvement. Bargaining is not automatically welfare-improving on top

Table 3: Consumer surplus share by seller information level.

Seller information level	Price	Consumer surplus	CS share
Low	615.3	378.6	0.84
Preference-aware	618.2	379.4	0.84
History-aware	623.0	376.5	0.83
High	696.4	304.5	0.68

Table 4: Buyer-side concentration (monopsony) effect on consumer surplus share.

Buyer-agent concentration	CS share
Common buyer-agent provider	0.80
Independent buyers	0.79

Table 5: Herfindahl-Hirschman index of seller win-shares, by regime and seller-agent concentration.

Regime	Independent sellers	Shared seller-agent provider
M0 (posted, limited search)	0.1015	0.1012
M1 (full-market search)	0.595	0.5499
M2 (bilateral, price-only)	0.595	0.5499
M3 (bilateral, multi-attribute)	0.595	0.5499
M4 (competitive)	0.4532	0.4128

Table 6: Seller-side agent-provider concentration: markup dispersion and price coefficient of variation.

Seller-agent concentration	Markup SD	Price CV
Independent sellers	0.068	0.060
Shared seller-agent provider	0.074	0.067

Table 7: Cluster-robust regression coefficients underlying the four central mechanisms (dependent variable and reference category noted in each row).

Mechanism	Coefficient (dependent variable; reference)	Coef.	p
Bargaining scope	Welfare; price-only vs. multi-attribute	-13.67	< 0.001
Regime (M1 vs. M0)	Welfare	26.67	< 0.001
Regime (M2 vs. M0)	Welfare	26.67	< 0.001
Seller information (high vs. low)	CS share	-0.160	< 0.001
Buyer concentration (independent vs. common provider)	CS share	-0.016	0.004
Seller concentration (shared vs. independent)	Price	-11.79	< 0.001
Seller information (high vs. low)	Price	81.10	< 0.001

Table 8: LLM-agent negotiation pilot: deal rate and welfare by regime and model pairing (N=36 episodes).

Regime	Model pairing (buyer-seller)	Deal rate	Mean price	Mean welfare
Price-only	Gemini-Gemini	100%	756.7	505.3
Price-only	Gemini-OpenRouter	83%	730.0	441.2
Price-only	OpenRouter-OpenRouter	100%	903.3	505.3
Multi-attribute	Gemini-Gemini	100%	769.2	539.5
Multi-attribute	Gemini-OpenRouter	67%	753.8	365.1
Multi-attribute	OpenRouter-OpenRouter	50%	873.3	271.5




Figure 1: Stacked consumer and producer surplus by institutional regime.




Figure 2: Consumer surplus share as a function of seller information about the buyer.




Figure 3: Welfare under price-only versus multi-attribute bargaining scope, by regime.




Figure 4: Herfindahl-Hirschman index of seller win-shares by regime and seller-agent concentration.

of that search; it is close to purely redistributive unless the bargaining space
605 is deliberately widened to include attributes beyond price, in which case it becomes a genuine source of new value. That distinction between search, price bargaining, and scope-expanding bargaining is exactly the kind of institutional granularity that the Coasean Singularity argument of Shahidi et al. (2025) anticipates in principle, and this paper supplies a quantitative decomposition of it
610 in one concrete setting.

The information result in Section 5.2 is the paper’s clearest illustration of why "agentic commerce" cannot be evaluated as a single undifferentiated technology from a consumer-welfare standpoint. The same AI agent that delivers the search-efficiency gains of Section 5.1 is, if it is also allowed to infer a buyer’s
615 urgency and purchase history, the mechanism by which a seller’s agent captures nearly a third of the surplus that a low-information seller would have left with the buyer. A platform, or a regulator, that wants the efficiency gains of agentic search without the extraction risk of agentic personalization cannot get both by simply deciding whether or not to "allow AI agents"; it has to decide,
620 specifically, what information a seller’s agent is permitted to infer and act on, a much more granular and more actionable design question than the current policy conversation about AI and e-commerce typically poses.

The concentration results in Sections 5.3 through 5.5 speak directly to the competition-policy literature this paper reviewed in Section 2.3, and to Bostoen
625 and Krämer (2026)’s question of whether existing regulation anticipates a world of autonomous negotiating agents. Search-driven concentration (Section 5.3) is a genuine, sizable effect – a roughly sixfold increase in the seller win-share Herfindahl index – but it is a consequence of efficient matching working as intended, not of any anticompetitive conduct, and the simulation shows a concrete, non-regulatory remedy already available to platform designers: requiring
630 or defaulting to competitive multi-seller solicitation (M4) meaningfully reverses much of that concentration without giving up the multi-attribute welfare gains of Section 5.1. The seller-infrastructure result (Section 5.5) requires more interpretive caution, precisely because it is the finding most likely to be over-read as

635 evidence of algorithmic collusion in the sense of Calvano et al. (2020) or Klein
(2021). This paper’s simulation contains no repeated-game punishment mechanism,
no signaling channel, and no learning dynamic of the kind those papers study;
the correlation this paper documents among sellers sharing a common pricing-agent
provider is a mechanical consequence of those sellers inheriting
640 correlated strategy parameters from their shared vendor, closer to a franchise
or platform-standardization effect than to tacit collusion. Regulators evaluating
whether a shared seller-agent vendor’s client network constitutes a competition
concern should, on the evidence here, look for the specific mechanism – shared
strategy parameters versus a genuine repeated-game equilibrium – rather than
645 inferring collusion from price correlation alone.

The buyer-side monopsony result (Section 5.4) is smaller in magnitude than
any of the above but is arguably the most forward-looking, since it is the one
result in this paper that scales with adoption in a way the others do not automatically:
if a small number of buyer-agent platforms eventually intermediate
650 a large share of retail transactions, the aggregated bargaining leverage documented
here at a 70-percent penetration rate inside one condition could plausibly grow,
and the natural extension – examining whether the monopsony effect is convex or
concave in buyer-agent-provider market share – is a well-posed empirical question
this paper’s simulation is equipped to answer in future work,
655 though it does not attempt to here.

7. Limitations

The central limitation of this paper is one it has tried to state honestly
throughout rather than discover in a closing section: the primary empirical engine
is a rule-based agent-based computational economics simulation, calibrated
660 qualitatively to the mechanisms described in the agentic-commerce literature
but not fit to, or validated against, any specific deployed large language model
or commercial AI shopping agent. The supplementary LLM-agent pilot in Sections
4.2 and 5.6 provides a small, genuinely encouraging piece of corroborating

evidence that the paper’s central bundling result is not an artifact of the rule-based model’s particular functional form, but thirty-six episodes across three
665 model pairings is an exploratory check, not a validation study, and should not be read as evidence that any specific commercial deployment will reproduce this paper’s numerical magnitudes.

Several further limitations follow from the simulation’s design choices. Buyer
670 valuations, seller costs, and bundle values are drawn from stylized, independent uniform distributions calibrated for tractability and interpretability rather than fit to any real product category’s actual price and cost data; the qualitative ordering of the results reported here (search improves matching, price-only bargaining redistributes, multi-attribute bargaining creates surplus) is a structural
675 property of the model that should be robust to reasonable changes in these distributions, but the specific magnitudes – the 13.6-unit welfare gain, the fall from 84 to 68 percent in consumer surplus share – are properties of this particular calibration and should not be read as forecasts for any specific real market. The simulation also studies a single, homogeneous product category; real e-commerce spans categories with very different search costs, negotiation norms,
680 and repeat-purchase dynamics, and nothing in this paper’s design speaks to how these results would generalize across categories. Finally, the paper works entirely with simulated rather than real transaction data, a consequence of the fact that autonomous buyer-seller agent negotiation, in the full sense studied here, is not
685 yet observable at the scale or with the institutional variation this paper’s design requires; validating these mechanism-level predictions against real transaction data, as agentic commerce platforms mature and such data becomes available, is the natural and necessary next step this paper’s own approach cannot take on its own.

690 8. Conclusion

This paper set out to ask whether the arrival of AI shopping and pricing agents matters for e-commerce markets because of what those agents are, or

because of the specific institutional rules under which they are allowed to search, bargain, and share information. The evidence assembled here, from a matched, fully reproducible simulation of 105,600 transactions and a smaller corroborating pilot with real language models, comes down clearly on the side of institutional rules. Search is close to automatically beneficial; price-only bargaining is close to purely redistributive; multi-attribute bargaining is where genuine efficiency gains are created; seller information is where those gains are put at risk of being extracted from the buyer who helped create them; and agent-provider concentration, on both sides of the market, leaves real but importantly different fingerprints depending on whether it is the buyer's side or the seller's side that is concentrated. None of these findings depends on any single AI system's particular training or prompting; all of them depend on design choices – how many sellers a shopping agent solicits, what a pricing agent is allowed to infer, how wide the negotiation space is allowed to be – that platform designers and regulators are making, or will soon be making, right now. The paper's central practical claim is that those design choices, not the underlying presence of AI, are where the welfare and competition consequences of agentic commerce will actually be decided.

CRedit authorship contribution statement

The author is solely responsible for all aspects of this work: conceptualization, methodology, software, validation, formal analysis, investigation, data curation, writing of the original draft, writing review and editing, and visualization.

Declaration of competing interest

The author declares no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Funding

720 This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Data availability

The simulated transaction dataset, the LLM-agent negotiation pilot’s episode data and transcripts, and all simulation, analysis, and figure-generation code
725 that support the findings of this study are openly available in a public code repository. A repository link will be provided upon acceptance to preserve the integrity of the double-anonymized review process.

Declaration of Generative AI and AI-assisted technologies in the writing process

730 During the preparation of this work, the author used Claude (Anthropic) to assist with the design and implementation of the agent-based simulation code, the statistical analysis and figure-generation scripts, the implementation of the supplementary LLM-agent negotiation pilot, and the drafting and editing of portions of this manuscript’s text. The author reviewed, verified, and edited
735 all AI-assisted content, independently confirmed all reported numerical results against the underlying code and data outputs, and takes full responsibility for the content of this publication.

References

- Acquisti, A., Varian, H.R., 2005. Conditioning prices on purchase history. Mar-
740 keting Science 24, 367–381.
- Akerlof, G.A., 1970. The market for "lemons": Quality uncertainty and the market mechanism. Quarterly Journal of Economics 84, 488–500.

- Allouah, A., Besbes, O., Figueroa, J.D., Kanoria, Y., Kumar, A., 2026. What is your ai agent buying? evaluation, biases, model dependence, & emerging implications of agentic e-commerce, in: Proceedings of the ACM Web Conference 2026 (WWW '26), pp. 8697–8700. doi:10.1145/3774904.3792943.
- Bakos, J.Y., 1997. Reducing buyer search costs: Implications for electronic marketplaces. *Management Science* 43, 1676–1692. doi:10.1287/mnsc.43.12.1676.
- Bansal, G., Hua, W., Huang, Z., Fournay, A., Swearngin, A., Epperson, W., Payne, T., Hofman, J.M., Lucier, B., Singh, C., et al., 2025. Magentic marketplace: An open-source environment for studying agentic markets. *ArXiv preprint arXiv:2510.25779*.
- Bergemann, D., Brooks, B., Morris, S., 2015. The limits of price discrimination. *American Economic Review* 105, 921–957.
- Bianchi, F., Chia, P.J., Yüsekşönül, M., Tagliabue, J., Jurafsky, D., Zou, J., 2024. How well can llms negotiate? negotiationarena platform and analysis, in: Proceedings of the 41st International Conference on Machine Learning (ICML 2024). *ArXiv:2402.05863*.
- Bichler, M., 2026. Agentic markets. *Electronic Markets* 36, 55. doi:10.1007/s12525-026-00906-y.
- Bichler, M., Durmann, J., Oberlechner, M., 2025. Agentic markets: Game dynamics and equilibrium in markets with learning agents. *ArXiv preprint arXiv:2506.18571*.
- Bostoen, F., Krämer, J., 2026. How future-proof is the dma? a case study of ai agents. *Journal of Competition Law & Economics* doi:10.1093/joclec/nhag005. advance article nhag005.
- Brown, Z.Y., MacKay, A., 2023. Competition in pricing algorithms. *American Economic Journal: Microeconomics* 15, 109–156.

- Calvano, E., Calzolari, G., Denicolò, V., Pastorello, S., 2020. Artificial intel-
770 ligence, algorithmic pricing, and collusion. *American Economic Review* 110,
3267–3297. doi:10.1257/aer.20190623.
- Calvano, E., Calzolari, G., Denicolò, V., Pastorello, S., 2021. Algorithmic col-
lusion with imperfect monitoring. *International Journal of Industrial Organi-
zation* 79, 102712. doi:10.1016/j.ijindorg.2021.102712.
- 775 Coase, R.H., 1937. The nature of the firm. *Economica* 4, 386–405.
- Dubé, J.P., Misra, S., 2023. Personalized pricing and consumer welfare. *Journal
of Political Economy* 131, 131–189.
- Faratin, P., Sierra, C., Jennings, N.R., 1998. Negotiation decision functions
for autonomous agents. *Robotics and Autonomous Systems* 24, 159–182.
780 doi:10.1016/S0921-8890(98)00029-3.
- Hadfield, G.K., Koh, A., 2025. An Economy of AI Agents. Technical Report.
National Bureau of Economic Research. NBER Working Paper; chapter in
The Economics of Transformative AI.
- Harrington, Joseph E., J., 2001. Comment on “reducing buyer search costs:
785 Implications for electronic marketplaces”. *Management Science* 47, 1727–
1732. doi:10.1287/mnsc.47.12.1727.10244.
- Kim, J., Kim, J.H., Kim, C., Park, J., 2023. Decisions with chatgpt: Reexam-
ining choice overload in chatgpt recommendations. *Journal of Retailing and
Consumer Services* 75, 103494. doi:10.1016/j.jretconser.2023.103494.
- 790 Klein, T., 2021. Autonomous algorithmic collusion: Q-learning under sequential
pricing. *RAND Journal of Economics* 52, 538–558.
- Nash, J.F., 1950. The bargaining problem. *Econometrica* 18, 155–162.
- Rubinstein, A., 1982. Perfect equilibrium in a bargaining model. *Econometrica*
50, 97–110. doi:10.2307/1912531.

- 795 Sanchez-Cartas, J.M., Katsamakas, E., 2025. Ai pricing algorithms under platform competition. *Electronic Commerce Research* 25, 4343–4370. doi:10.1007/s10660-024-09821-w.
- Shahidi, P., Rusak, G., Manning, B.S., Fradkin, A., Horton, J.J., 2025. The Coasean Singularity? Demand, Supply, and Market Design with AI Agents. Technical Report 34468. National Bureau of Economic Research. NBER
800 Working Paper 34468; chapter 6 in *The Economics of Transformative AI*, University of Chicago Press.
- Stigler, G.J., 1961. The economics of information. *Journal of Political Economy* 69, 213–225.
- 805 Tesfatsion, L., 2006. Agent-based computational economics: A constructive approach to economic theory, in: Tesfatsion, L., Judd, K.L. (Eds.), *Handbook of Computational Economics, Volume 2: Agent-Based Computational Economics*. Elsevier/North-Holland, pp. 831–880. doi:10.1016/S1574-0021(05)02016-0.
- 810 Williamson, O.E., 1981. The economics of organization: The transaction cost approach. *American Journal of Sociology* 87, 548–577.
- Zhu, S., Sun, J., Nian, Y., South, T., Pentland, A., Pei, J., 2025. The automated but risky game: Modeling and benchmarking agent-to-agent negotiations and transactions in consumer markets, in: *Proceedings of the Natural Language Processing Workshop (NLLP 2025)*. ACL Anthology 2025.nllp-1.2; also
815 arXiv:2506.00073.